


Wrangle Data


Sometimes the data we are working with is not quite ready for whatever it is we want to do.



`dplyr`: your new best friend.

- Three key functions:
 - `filter()` to subset data based on rows






`dplyr`

- Three key functions:
 - `filter()` to subset data based on rows
 - `select()` to subset data based on columns



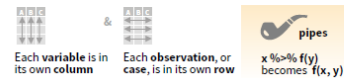
`dplyr`

- Three key functions:
 - ``filter()`` to subset data based on rows
 - ``select()`` to subset data based on columns
 - ``mutate()`` to add or modify values in columns




Data Transformation with dplyr : : CHEAT SHEET

dplyr functions work with pipes and expect tidy data. In tidy data:



Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).




VARIATIONS

summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.

Group Cases

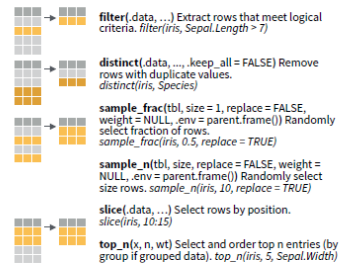
Use **group_by**() to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.



Manipulate Cases

EXTRACT CASES

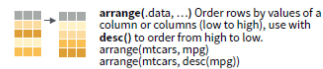
Row functions return a subset of rows as a new table.



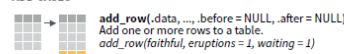
Logical and boolean operators to use with filter()

< <= is.na() %in% | xor()
> >= is.na() ! &
See ?base::Logic and ?Comparison for help.

ARRANGE CASES



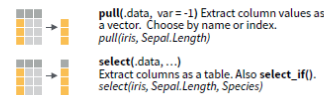
ADD CASES



Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.



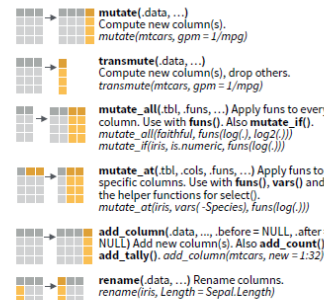
Use these helpers with **select** (), e.g. *select(iris, starts_with("Sepal"))*

contains(match) **num_range**(prefix, range) , e.g. *mpg:cyl*
ends_with(match) **one_of**(...) , e.g. *Species*
matches(match) **starts_with**(match)

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

vectorized function



Vector Functions

TO USE WITH MUTATE ()

mutate() and **transmute**() apply vectorized functions to columns to create new columns. Vectorized functions take vectors as input and return single vectors of the same length as output.

vectorized function

OFFSETS

dplyr::lag() - Offset elements by 1
dplyr::lead() - Offset elements by -1

CUMULATIVE AGGREGATES

dplyr::cumall() - Cumulative all()
dplyr::cumany() - Cumulative any()
dplyr::cummax() - Cumulative max()
dplyr::cummean() - Cumulative mean()
dplyr::cummin() - Cumulative min()
dplyr::cumprod() - Cumulative prod()
dplyr::cumsum() - Cumulative sum()

RANKINGS

dplyr::cume_dist() - Proportion of all values <= *dplyr::dense_rank()* - rank w ties = min, no gaps
dplyr::min_rank() - rank with ties = min
dplyr::ntile() - bins into n bins
dplyr::percent_rank() - min_rank scaled to [0,1]
dplyr::row_number() - rank with ties = "first"

MATH

+, *****, **/**, **^**, **%%**, **%/%**, **%/%** - arithmetic ops
log(), **log2()**, **log10()** - logs
<, **<=**, **>**, **>=**, **is**, **==** - logical comparisons
dplyr::between() - *x >= left & x <= right*
dplyr::near() - *safe ==* for floating point numbers

MISC

dplyr::case_when() - multi-case if_else()
iris %>% mutate(Species = case_when(Species == "versicolor" ~ "vers", Species == "virginica" ~ "virgi", TRUE ~ Species))

dplyr::coalesce() - first non-NA values by element across a set of vectors
dplyr::if_else() - element-wise if() + else()
dplyr::na_if() - replace specific values with NA
pmax() - element-wise max()
pmin() - element-wise min()
dplyr::recode() - Vectorized switch()
dplyr::recode_factor() - Vectorized switch() for factors

Summary Functions

TO USE WITH SUMMARISE ()

summarise() applies summary functions to columns to create a new table. Summary functions take vectors as input and return single values as output.

summary function

COUNTS

dplyr::n() - number of values/rows
dplyr::distinct() - # of uniques
sum(is.na()) - # of non-NA's

LOCATION

mean() - mean, also **mean(is.na())**
median() - median

LOGICALS

sum() - Proportion of TRUE's
mean() - # of TRUE's

POSITION/ORDER

dplyr::first() - first value
dplyr::last() - last value
dplyr::nth() - value in nth location of vector

RANK


quantile() - nth quantile
min() - minimum value
max() - maximum value

SPREAD

IQR() - Inter-Quartile Range
mad() - median absolute deviation
sd() - standard deviation
var() - variance

Row Names


Tidy data does not use rownames, which store a variable outside of the columns. To work with the rownames, first move them into a column.



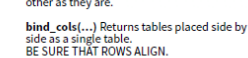
Also has **rownames()**, **remove_rownames()**

Combine Tables


COMBINE VARIABLES




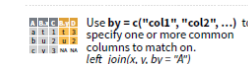
COMBINE CASES



Use a "Mutating Join" to join one table to columns from another, matching values with the rows that they correspond to. Each join retains a different combination of values from the tables.



EXTRACT ROWS

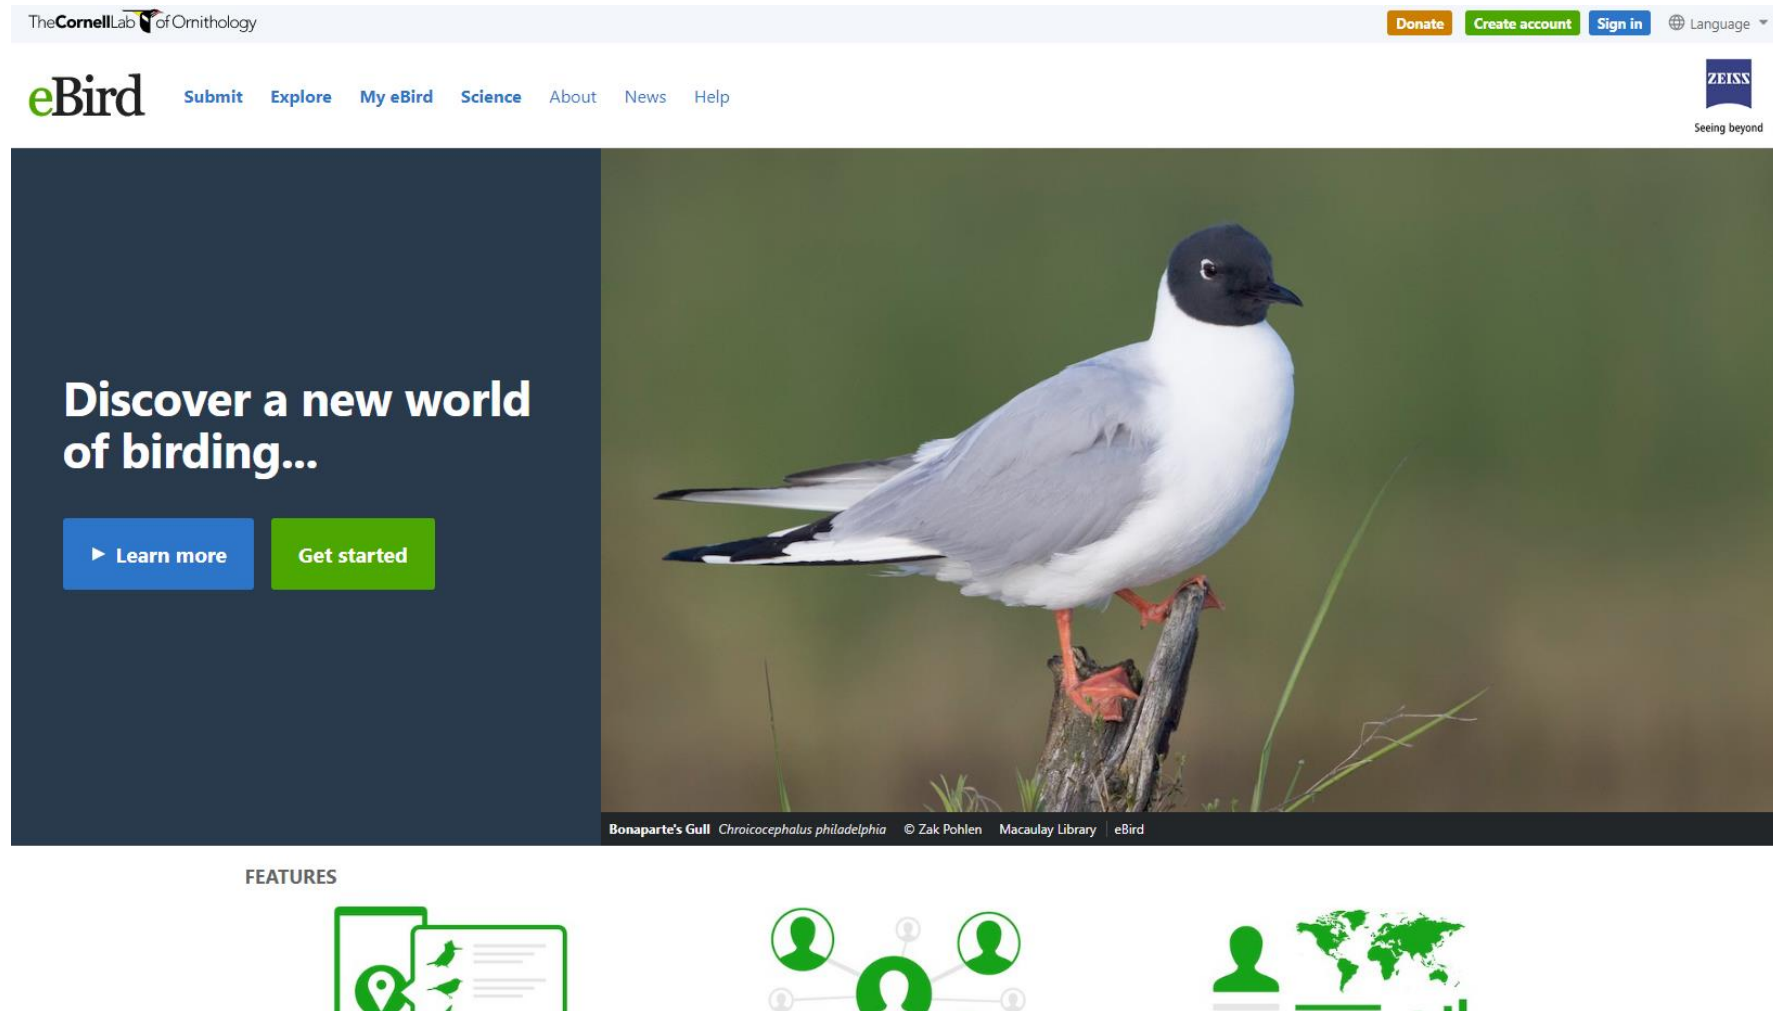


RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with brows/vignettes/package = c("dplyr", "tidbtl") • dplyr 0.7.0 • tidbtl 1.2.0 • Updated: 2019-08



RStudio® is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more with brows/vignettes/package = c("dplyr", "tidbtl") • dplyr 0.7.0 • tidbtl 1.2.0 • Updated: 2019-08

About the data



About the data

```
ebird <- read.csv(here::here("data",  
"eBird_workshop.csv", stringsAsFactors = FALSE))
```

Don't forget to check your data!

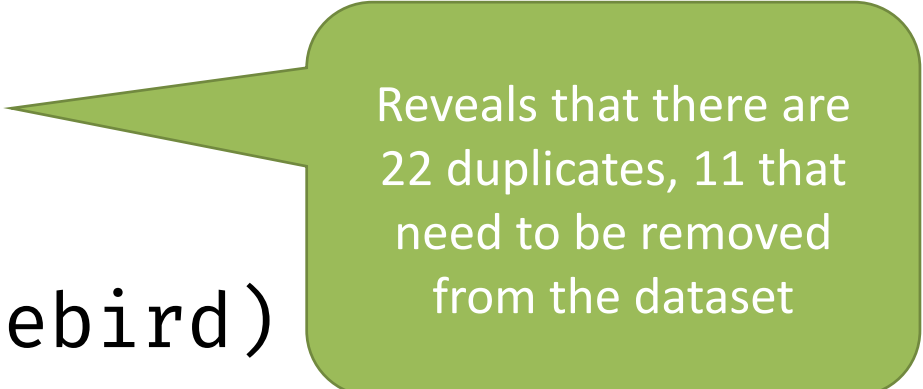
```
dplyr::glimpse(ebird)
```


About the data


There are duplicate rows in this data frame...but! we can easily figure that out.

```
janitor::get_dupes(ebird)
```

```
ebird <- dplyr::distinct(ebird)  
dplyr::glimpse(ebird)
```



Reveals that there are 22 duplicates, 11 that need to be removed from the dataset



`filter()`

- The eBird data set has a lot of rows. Let's examine entries from a particular state.

```
filter(ebird, state == 'AK')
```

```
ebird %>%  
  filter(state == 'AK')
```

We tend to use
pipes!

$x \%>\% f(y) = f(x,y)$
"and then..."

`filter()`

- We use `==` to specify an exact condition
 - `<` , less than
 - `<=` , less than or equal to
 - `==` , equals
 - `!=` , *not* equal to
 - `>=` , greater than or equal to
 - `>` , greater than
- The condition must return either true or false, for each row.



See `?base::Logic`
and `?Comparison`
for help

`filter()` with multiple conditions

You can filter on more than one criterion

- Example: we want all the birds from Alaska, but only in the year 2008

```
ebird %>%  
  filter(state == 'AK',  
         year == 2008)
```

It does not matter
if you select
`state` or
`year` first, you
can do either.

`filter()` with multiple conditions

You can filter on more than one criterion

- What if we want to look at birds from more than one state?

```
my_states <- c('AK', 'AL', 'MS')
```

```
ebird %>%
```

```
  filter(year == 2008,  
         state %in% my_states  
        )
```



x %in% y

``filter()`` with multiple conditions

You can filter on more than one criterion

- What if we want to look at birds from more than one state?

```
ebird %>%  
  filter(year == 2008,  
         state %in% c('AK', 'AL', 'MS'))
```

`filter()` with multiple conditions

You can filter on more than one criterion

- Multiple arguments to `filter()` are combined with “and”:






Figure 5.1: Complete set of boolean operations. x is the left-hand circle, y is the right-hand circle, and the shaded region show which parts each operator selects.

YOUR TURN 1

1. How could we pull out birds in Alaska (AK), before 2010? *(Hint: you can use the same symbols on year that you would use with any other numbers)*
2. Filter the data to contain only the species "American Coot" from MS and FL (your instructors' states), in all years **except** 2010. *Assign this object to a data frame and verify (using ``unique()``) that you did it right.*

YOUR TURN: how did you do?





`select()`

- Simplify your data by keeping only the columns you want to work with. You can do this two ways:

```
ebird %>%  
  select(species, state, year)
```

```
ebird %>%  
  select(-samplesize, -presence)
```

``select()`` order DOES matter

You can rearrange your columns of data this way

```
ebird %>%  
  select(species, state, year)
```

species	state	year
American Coot	Florida	2008

`select()` order DOES matter

You can rearrange your columns of data this way

```
ebird %>%  
  select(year, species, state)
```

year	species	state
2008	American Coot	Florida

`select()` order DOES matter

You can also just move one or a few column and keep the original order of all the other columns

```
ebird %>%  
  select(year, everything())
```

`select()` has helper functions

`select(ebird, starts_with("s"))`

Data Transformation with dplyr :: CHEAT SHEET

dplyr functions work with pipes and expect tidy data. In tidy data:

Each variable is in its own column
Each observation, or case, is in its own row
pipes
x %>% f(y) becomes f(x, y)

Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

summary function
`summarise(.data, ...)`
Compute table of summaries.
`summarise(mtcars, avg = mean(mpg))`
`count(x, ..., wt = NULL, sort = FALSE)`
Count number of rows in each group defined by the variables in ... Also `tally()`.
`count(iris, Species)`

VARIATIONS

`summarise_all()` - Apply funs to every column.
`summarise_at()` - Apply funs to specific columns.
`summarise_if()` - Apply funs to all cols of one type.

Group Cases

Use `group_by()` to create a "grouped" copy of a table. dplyr functions will manipulate each "group" separately and then combine the results.

`mtcars %>%
group_by(cyl) %>%
summarise(avg = mean(mpg))`

group_by(.data, ..., add = FALSE)

Returns copy of table grouped by ...
`g_iris <- group_by(iris, Species)`

Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.

`filter(.data, ...)` Extract rows that meet logical criteria. `filter(iris, Sepal.Length > 7)`
`distinct(.data, ..., keep_all = FALSE)` Remove rows with duplicate values.
`distinct(iris, Species)`
`sample_frac(tbl, size = 1, replace = FALSE, weight = NULL, env = parent.frame())` Randomly select fraction of rows.
`sample_frac(iris, 0.5, replace = TRUE)`
`sample_n(tbl, size, replace = FALSE, weight = NULL, env = parent.frame())` Randomly select size rows.
`sample_n(iris, 10, replace = TRUE)`
`slice(.data, ...)` Select rows by position.
`slice(iris, 10:15)`
`top_n(x, n, wt)` Select and order top n entries (by group if grouped data).
`top_n(iris, 5, Sepal.Width)`

Logical and boolean operators to use with filter()

< <= is.na() %in% | xor()
> >= is.na() %in% |
See ?base::Logic and ?Comparison for help.

ARRANGE CASES

`arrange(.data, ...)` Order rows by values of a column or columns (low to high), use with `desc()` to order from high to low.
`arrange(mtcars, mpg)`
`arrange(mtcars, desc(mpg))`

ADD CASES

`add_row(.data, ..., before = NULL, after = NULL)`
Add one or more rows to a table.
`add_row(faithful, eruptions = 1, waiting = 1)`
`add_rows(.data, rows)`
Add multiple rows to a table.
`add_rows(faithful, eruptions = 1:10, waiting = 1:10)`

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.

`pull(.data, var = -1)` Extract column values as a vector. Choose by name or index.
`pull(iris, Sepal.Length)`
`select(.data, ...)` Extract columns as a table. Also `select_if()`.
`select(iris, Sepal.Length, Species)`
Use these helpers with `select()`, e.g. `select(iris, starts_with("Sepal"))`
`contains(match)` `num_range(prefix, range)` ; e.g. `mpg:cyl`
`ends_with(match)` `one_of(...)` ; e.g. `-Species`
`matches(match)` `starts_with(match)`

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).

vectorized function
`mutate(.data, ...)` Compute new column(s).
`mutate(mtcars, gpm = 1/mpg)`
`transmute(.data, ...)` Compute new column(s), drop others.
`transmute(mtcars, gpm = 1/mpg)`
`mutate_all(.tbl, funs, ...)` Apply funs to every column. Use with `funs()`. Also `mutate_at()`.
`mutate_all(faithful, funs(log(), log2()))`
`mutate_if(iris, is.numeric, funs(log()))`
`mutate_at(.tbl, cols, funs, ...)` Apply funs to specific columns. Use with `funs()`, `vars()` and the helper functions for `select()`.
`mutate_at(iris, vars(Species), funs(log()))`
`add_column(.data, ..., before = NULL, after = NULL)` Add new column(s). Also `add_count()`.
`add_tally()` `add_column(mtcars, new = 1:32)`
`rename(.data, ...)` Rename columns.
`rename(iris, Length = Sepal.Length)`

Manipulate Variables

EXTRACT VARIABLES

Column functions return a set of columns as a new vector or table.

`pull(.data, var = -1)` Extract column values as a vector. Choose by name or index.
`pull(iris, Sepal.Length)`

`select(.data, ...)`
Extract columns as a table. Also `select_if()`.
`select(iris, Sepal.Length, Species)`


Use these helpers with `select()`,
e.g. `select(iris, starts_with("Sepal"))`

`contains(match)` `num_range(prefix, range)` ; e.g. `mpg:cyl`
`ends_with(match)` `one_of(...)` ; e.g. `-Species`
`matches(match)` `starts_with(match)`



RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 848-448-1212 • rstudio.com • Learn more with browser/extensions/package = "dplyr", "tidyr" • dplyr 0.7.0 • tidyr 1.2.0 • Updated: 2019-08


YOUR TURN 2



From the ebird data, subset to only include the species American Coot, from the states FL, AL, and MS. Keep only the `state`, `year`, and `presence` columns.

What is the proper order of operations in this case?

YOUR TURN: how did you do?




About the data

Daily averages of water quality data (e.g., temperature, salinity, dissolved oxygen, etc.) from 6 National Estuarine Research Reserves

```
```{r}
wq <- read.csv(here::here("data", "daily_wq.csv"), stringsAsFactors = FALSE)
glimpse(wq)
```
```

YOUR TURN 3



We have read in the daily water quality data for a few stations. Create a new data frame called `wq_trimmed` where, from `wq`, you:

1. **Select** the following columns: `station_code`, `month`, `day`, `temp`, `sal`, `do_pct`, and `depth`.
2. **Filter** for rows where `depth` is *not* missing. (Hint: `is.na` is the function that checks to see if a value *is* missing. How would you look for "not" `is.na`? It's similar to "not equal to" ...)

YOUR TURN: how did you do?




How much was removed from the data frame?



Modifying with `mutate()`





Modifying with `mutate()`

- `mutate()` operates on **rows**
- For example, depth is recorded in meters, perhaps feet would be better for the public in the US:

```
wq_trimmed %>%  
  mutate(depth_ft = depth * 3.28)
```

Why 3.28?!

Modifying with `mutate()`

- You can also use other columns AND use a column after creating it



Modifying with `mutate()`

- You can also use other columns AND use a column after creating it

```
wq_trimmed <- wq_trimmed %>%  
  mutate(monthday = paste(month, day, sep = "-"),  
         meaningless = sal + temp,  
         even_more_meaningless = meaningless + 5)
```


YOUR TURN 4

There are two parts to this. You can approach them separately or within the same series of pipes. Remember to save the result as the new, better, `wq_trimmed` data frame!

1. Remove `monthday` and the `meaningless_things` from the `wq_trimmed` data frame.
2. The same person that wants to see `depth` in feet rather than meters *also* wants you to turn `temp` into Fahrenheit, from Celsius. You've looked up the conversion. Now create a new column, `temp_f`, with the new variable.
`F = (9/5)(temp in C) + 32**`**

YOUR TURN: how did you do?



Group-wise operations with `group_by()` and `summarize()`

- `group_by()` changes each function from operating on the full dataset to specified groups. *This can be done in conjunction with other `dplyr` functions!*
- `summarize()` reduces multiple values down to a single summary

```
ebird %>%
```

```
  group_by(state) %>%
```

```
  summarize(mean_presence = mean(presence, na.rm = TRUE),  
            max_presence = max(presence, na.rm = TRUE)  
            )
```

Group-wise operations with `group_by()` and `summarize()`

- `group_by()` changes each function from operating on the full dataset to specified groups. *This can be done in conjunction with other `dplyr` functions!*
- `summarize()` reduces multiple values down to a single summary

```
ebird %>%
```

```
  group_by(state, species) %>%
```

```
  summarize(mean_presence = mean(presence, na.rm = TRUE),  
            max_presence = max(presence, na.rm = TRUE)  
            )
```

YOUR TURN 5

Group the `wq_trimmed` dataset to calculate monthly average temperature and salinity, and their standard deviations, at each station.

YOUR TURN: how did you do?



Variations on `summarize()``

Summarise Cases

These apply **summary functions** to columns to create a new table of summary statistics. Summary functions take vectors as input and return one value (see back).

summary function



summarise(.data, ...)
Compute table of summaries.
summarise(mtcars, avg = mean(mpg))




count(x, ..., wt = NULL, sort = FALSE)
Count number of rows in each group defined by the variables in ... Also **tally()**.
count(iris, Species)

VARIATIONS

summarise_all() - Apply funs to every column.
summarise_at() - Apply funs to specific columns.
summarise_if() - Apply funs to all cols of one type.

Data Transformation with dplyr : : CHEAT SHEET

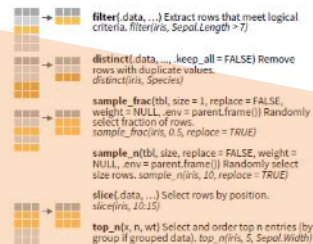
dplyr functions work with pipes and expect tidy data. In tidy data:



Manipulate Cases

EXTRACT CASES

Row functions return a subset of rows as a new table.



Logical and boolean operators to use with filter()

< <= is.na() %in% | xor()
> >= !is.na() ! &
See **Base::Logic** and **TComparison** for help.

ARRANGE CASES

arrange(.data, ...) Order rows by values of a column or columns (low to high), use with **desc()** to order from high to low.
arrange(mtcars, mpg)
arrange(mtcars, desc(mpg))


ADD CASES

add_row(.data, ..., before = NULL, after = NULL)
Add one or more rows to a table.
add_row(faithful, eruptions = 1, waiting = 1)

Manipulate Variables

EXTRACT VARIABLES

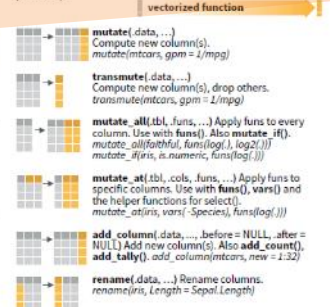
Column functions return a set of columns as a new vector or table.



Use these helpers with **select()**,
e.g. *select(iris, starts_with("Sepal"))*
contains(match) **num_range(prefix, range)** t, e.g. *mpg:cyl*
ends_with(match) **one_of(...)** e.g. *Species*
matches(match) **starts_with(match)**

MAKE NEW VARIABLES

These apply **vectorized functions** to columns. Vectorized funs take vectors as input and return vectors of the same length as output (see back).



RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio - info@rstudio.com • +44-448-1212-rstudio.com • Learn more with browser/printed/package = c("dplyr", "RStudio") • dplyr: 0.7.0 • RStudio: 1.2.0 • Updated: 2019-08



Sorting your data with `arrange()`

Let's put our ebird summary in order by species and then state.

Or put it in order by max presence



YOUR TURN 6

Use the `wq_trimmed` data frame. Calculate monthly average temp, sal, and do_pct (at least - more variables if you like) *for each station*.

Make a scatterplot using any two of these variables. Use what you've learned about `ggplot2`'s options to adjust the look and feel of the graph.

Is the relationship what you expected it to be?

Does it vary by site?


YOUR TURN: How did you do?



"Data comes in many formats, but R prefers just one: tidy data. "

- Garrett Grolemund

Dates and Times



Dates and times with lubridate : : CHEAT SHEET

Date-times



2017-11-28 12:00:00
A date-time is a point on the timeline, stored as the number of seconds since 1970-01-01 00:00:00 UTC
dt <- as_datetime(1511870400)
"2017-11-28 12:00:00 UTC"

2017-11-28
A date is a day stored as the number of days since 1970-01-01
d <- as_date(17498)
"2017-11-28"

12:00:00
An hms is a time stored as the number of seconds since 00:00:00
t <- hms::as_hms(85)
"00:01:25"

PARSE DATE-TIMES (Convert strings or numbers to date-times)
1. Identify the order of the year (y), month (m), day (d), hour (h), minute (m) and second (s) elements in your data
2. Use the function below whose name replicates the order. Each accepts a wide variety of input formats.

2017-11-28T14:02:00 ymd_hms(), ymd_hm(), ymd_h(), ymd_hms("2017-11-28T14:02:00")

2017-22-12 10:00:00 ydm_hms(), ydm_hm(), ydm_h(), ydm_hms("2017-22-12 10:00:00")

11/28/2017 1:02:03 mdy_hms(), mdy_hm(), mdy_h(), mdy_hms("11/28/2017 1:02:03")

1 Jan 2017 23:59:59 dmy_hms(), dmy_hm(), dmy_h(), dmy_hms("1 Jan 2017 23:59:59")

20170131 ymd(), ydm(), ymd(20170131)

July 4th, 2000 mdy(), mdy(), mdy("July 4th, 2000")

4th of July 99 dmy(), dym(), dmy("4th of July 99")

2001:Q3 yq(), yq(), yq("2001:Q3")

2:01 hms::hms() Also lubridate::hms(), hm() and ms(), which return periods. hms::hms(sec = 0, min = 1, hours = 2)

2017.5 date_decimal(decimal, tz = "UTC") Q for quarter, date_decimal(2017.5)

now(tzone = "Z") Current time in tz (defaults to system tz), now()

today(tzone = "Z") Current date in a tz (defaults to system tz), today()

fast_strptime() Faster strptime, fast_strptime("9/1/01", "%y/%m/%d")

parse_date_time() Easier strptime, parse_date_time("9/1/01", "ymd")

GET AND SET COMPONENTS
Use an accessor function to get a component. Assign into an accessor function to change a component in place.
d <- "2017-11-28"
day(d) ## 28
day(d) <- 1
d ## "2017-11-01"

2018-01-31 11:59:59 date(x) Date component, date(dt)

2018-01-31 11:59:59 year(x) Year, year(dt)

2018-01-31 11:59:59 isoyear(x) The ISO 8601 year.

2018-01-31 11:59:59 epiyear(x) Epidemiological year.

2018-01-31 11:59:59 month(x, label, abbr) Month, month(dt)

2018-01-31 11:59:59 day(x) Day of month, day(dt)

2018-01-31 11:59:59 wday(x, label, abbr) Day of week, qday(x) Day of quarter.

2018-01-31 11:59:59 hour(x) Hour, hour(dt)

2018-01-31 11:59:59 minute(x) Minutes, minute(dt)

2018-01-31 11:59:59 second(x) Seconds, second(dt)

week(x) Week of the year, week(dt)

isoweek(x) ISO 8601 week.

epiweek(x) Epidemiological week.

quarter(x, with_year = FALSE) Quarter, quarter(dt)

semester(x, with_year = FALSE) Semester, semester(dt)

am(x) Is it in the am? am(dt)

pm(x) Is it in the pm? pm(dt)

dst(x) Is it daylight savings? dst(dt)

leap_year(x) Is it a leap year? leap_year(dt)

update(object, ..., simple = FALSE) update(dt, mday = 2, hour = 1)

Round Date-times



floor_date(x, unit = "second") Round down to nearest unit.

floor_date(dt, unit = "month") Round down to nearest unit.

round_date(x, unit = "second") Round to nearest unit.

round_date(dt, unit = "month") Round to nearest unit.

ceiling_date(x, unit = "second") Round up to nearest unit.

ceiling_date(dt, unit = "month") Round up to nearest unit.

rollback(dates, roll, to, first = FALSE, preserve_hms = TRUE) Roll back to last day of previous month, rollback(dt)

Stamp Date-times

stamp() Derive a template from an example string and return a new function that will apply the template to date-times. Also stamp_date() and stamp_time().

1. Derive a template, create a function
sf <- stamp("Created Sunday, Jan 17, 1999 3:34")
2. Apply the template to dates
sfymd("2010-04-03")
[1] "Created Monday, Apr 05, 2010 00:00"

Tip: use a date with day = 12

Time Zones

R recognizes ~600 time zones. Each encodes the time zone, Daylight Savings Time, and historical calendar variations for an area. R assigns one time zone per vector.

Use the UTC time zone to avoid Daylight Savings.

olsonNames() Returns a list of valid time zone names. OlsonNames()

master(x, with_year = FALSE) Master, master(dt)

am(x) Is it in the am? am(dt)

pm(x) Is it in the pm? pm(dt)

dst(x) Is it daylight savings? dst(dt)

leap_year(x) Is it a leap year? leap_year(dt)

update(object, ..., simple = FALSE) update(dt, mday = 2, hour = 1)

force_tz(time, tz = "Z") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

force_tz(dt, "US/Pacific") Get the same date-time in a new time zone (a new date-time).

Math with Date-times

Math with date-times relies on the timeline, which behaves inconsistently. Consider how the timeline behaves during:

A normal day
nor <- ymd_hms("2018-01-01 01:30:00", tz = "US/Eastern")

The start of daylight savings (spring forward)
gap <- ymd_hms("2018-03-11 01:30:00", tz = "US/Eastern")

The end of daylight savings (fall back)
lap <- ymd_hms("2018-11-04 00:30:00", tz = "US/Eastern")

Leap years and leap seconds
leap <- ymd("2019-03-01")

Periods track changes in clock times, which ignore time line irregularities.

normal + minutes(90)

gap + minutes(90)

lap + minutes(90)

leap + years(1)

Durations track the passage of physical time, which deviates from clock time when irregularities occur.

normal + dminutes(90)

gap + dminutes(90)

lap + dminutes(90)

leap + dyears(1)

Intervals represent specific intervals of the timeline, bounded by start and end date-times.

interval(normal, normal + minutes(90))

interval(gap, gap + minutes(90))

interval(lap, lap + minutes(90))

interval(leap, leap + years(1))

Not all years are 365 days due to leap days.

Not all minutes are 60 seconds due to leap seconds.

It is possible to create an imaginary date by adding months, e.g. February 31st

jan31 <- ymd(20180131)
jan31 + months(1)
NA

%m-% and %m-% will roll imaginary dates to the last day of the previous month.

jan31 %m-% months(1)
"2018-02-28"

add_with_rollback(e1, e2, roll_to_first = TRUE) will roll imaginary dates to the first day of the new month.

add_with_rollback(jan31, months(1), roll_to_first = TRUE)
"2018-03-01"

PERIODS

Add or subtract periods to model events that happen at specific clock times, like the NYSE opening bell.

Make a period with the name of a time unit pluralized, e.g.

p <- months(3) + days(12)
p
"3m 12d 0H 0M 0S"

years(x = 1) x years.
months(x = 1) x months.
weeks(x = 1) x weeks.
days(x = 1) x days.
hours(x = 1) x hours.
minutes(x = 1) x minutes.
seconds(x = 1) x seconds.
milliseconds(x = 1) x milliseconds.
microseconds(x = 1) x microseconds.
nanoseconds(x = 1) x nanoseconds.
picoseconds(x = 1) x picoseconds.

period(num = NULL, units = "second", ...) An automation friendly period constructor. period(5, unit = "years")

as.period(x, unit) Coerce a timespan to a period, optionally in the specified units. Also is.period(), as.period()

period_to_seconds(x) Convert a period to the "standard" number of seconds implied by the period. Also seconds_to_period(), period_to_seconds(p)

DURATIONS

Add or subtract durations to model physical processes, like battery life. Durations are stored as seconds, the only time unit with a consistent length. DiffTimes are a class of durations found in base R.

Make a duration with the name of a period prefixed with a d, e.g.

dd <- ddays(14)
dd
"1209600s (~2 weeks)"

Exact length in seconds

Equivalent in common units

ddays(x = 1) 31536000 x seconds.
dweeks(x = 1) 604800 x seconds.
dhours(x = 1) 3600 x seconds.
dminutes(x = 1) 60 x seconds.
dseconds(x = 1) x seconds.
dmilliseconds(x = 1) x 10⁻³ seconds.
dmicroseconds(x = 1) x 10⁻⁶ seconds.
dnanoseconds(x = 1) x 10⁻⁹ seconds.
dpicoseconds(x = 1) x 10⁻¹² seconds.

duration(num = NULL, units = "second", ...) An automation friendly duration constructor. duration(5, unit = "years")

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()

as.duration(x, ...) Coerce a timespan to a duration. Also is.duration(), as.duration()



RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more at lubridate.tidyverse.org • lubridate 1.6.0 • Updated: 2017-12



RStudio is a trademark of RStudio, Inc. • CC BY SA RStudio • info@rstudio.com • 844-448-1212 • rstudio.com • Learn more at lubridate.tidyverse.org • lubridate 1.6.0 • Updated: 2017-12

YOUR TURN

``mutate`` can be used to add a complete vector of the same value.

Fill in the skeleton code to create a column for a full yyyy-mm-dd style date, then make a **line** graph of temperature throughout the year at the water quality stations, and color them by `station_code`.

Which of these stations do you think is at the NERR in Alaska?